

# LoopNote — paid social creative

Newform take-home · Avi Kodwani

## What the data says

The account is at **\$105.12 per paid start against a \$90 target**, and six of ten tests missed. The useful finding isn't which ones lost — it's that **the four ads with the best attention and the four with the best trial → paid conversion share zero rows**. C10 posted the best thumbstop in the test (50%) and the worst conversion (11.5%, \$190 CPA). C05 posted 39% and converted at 21.0% for **\$61.97**. Attention isn't the constraint; it's being bought from the wrong people. The mechanism underneath is that **the ad sets the expectation the trial has to satisfy** — the skit promised a joke and delivered software (14.0% paid), the demo promised exactly what the app does (21.0%), the before/after promised a transformation (26.8%, best in the test). Which yields one rule: **show the transformation, don't describe it**. The five ads where the product's output is on screen blended **\$76.61**; the five where someone talks about it blended **\$192.71** — at nearly identical average thumbstop (36% vs 34%). The entire difference happens after the click.

| COHORT                     | ROWS                    | SPEND    | PAID | CPA             | TRIAL → PAID |
|----------------------------|-------------------------|----------|------|-----------------|--------------|
| Product's output on screen | C01, C02, C05, C07, C09 | \$22,600 | 295  | <b>\$76.61</b>  | <b>23.1%</b> |
| Someone describes it       | C03, C04, C06, C08, C10 | \$18,500 | 96   | <b>\$192.71</b> | <b>13.6%</b> |

Two supporting reads. **C03 is not a message failure** — the founder talking head converts trial → paid at 22.2%, fourth-best of ten, on the worst thumbstop (18%) in the test. Its argument works; its container is invisible. And **TikTok is bimodal, not weak**: it holds the best row in the account (C05, \$61.97) and three of the five worst. It punishes entertainment framing and rewards demo framing harder than Meta does.

Every script below opens on the mess rather than the clean version, and closes on a number that belongs to the person holding the phone rather than an industry statistic.

## Script A — "Thirty-one uhhs"

TikTok-first · 9:16 · 24s · lineage **C05** (\$61.97, best row) + C09's numeric hook (41% thumbstop)

| TIME      | ON SCREEN                                                                                                     | VO                                                                               | SOUND                   |
|-----------|---------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|-------------------------|
| 0:00–0:03 | Transcript screen. Search field on <b>uhh</b> . Counter reads <b>31</b> , highlights down the page.           | "My transcript says 'uhh' thirty-one times."                                     | Room tone. No music.    |
| 0:03–0:07 | Thumb scrolls fast. Legible: "...the thing with Priya — no wait, that was Tuesday — okay so if we push it..." | "That's four minutes of me in a parking garage, talking to myself after a call." | Room tone.              |
| 0:07–0:10 | Still scrolling. Scrollbar shows how much is left.                                                            | "I've never played one of these back. Not once."                                 | Room tone.              |
| 0:10–0:12 | One tap. Processing.                                                                                          | (silence)                                                                        | UI tick. Music in, low. |
| 0:12–0:18 | Task list resolves. Four items, two with names, two overdue. <b>Held long enough to read.</b>                 | "Four tasks. Two with names on them. One I'd completely forgotten I said."       | Music low.              |
| 0:18–0:21 | Split: the messy line ↔ the task it produced.                                                                 | "That's the raw. That's what it got out of it."                                  | Music.                  |
| 0:21–0:24 | LoopNote mark. On-screen: <b>Paid ad. Real memo .</b>                                                         | "LoopNote made this ad. The memo's mine. <b>Link's in the caption.</b> "         | Music out.              |

**Alt hooks:** "Read that and tell me a computer could do anything with it." · "This is four minutes of me talking to myself in a parking garage." **Proof:** one unbroken screen recording; raw input and output side by side at 0:18 so the viewer grades the extraction. **Objection answered:** accuracy — structurally. Nobody says it's accurate; you read it and decide. Saying it would break the rule the data gave us. **Predicted:** thumbstop 37–40%, trial → paid 22–25%, **CPA at or under \$61.97.**

---

## Script B — "Sixty-three memos"

---

**Meta-first** · 9:16 + 4:5 · 25s · lineage **C09** (\$82.09) + **C01** (\$81.36), the two largest budgets in the test

| TIME          | ON SCREEN                                                                     | VO                                                           | SOUND          |
|---------------|-------------------------------------------------------------------------------|--------------------------------------------------------------|----------------|
| 0:00–<br>0:04 | Phone over a kitchen table. Voice Memos list scrolling — long. Count visible. | "I had sixty-three voice memos I'd never played back."       | Room tone.     |
| 0:04–<br>0:07 | Bottom of the list. Dates from last year.                                     | "Some of them are a year old."                               | Room tone.     |
| 0:07–<br>0:10 | Face, briefly. Unstyled, morning light.                                       | "So on a Sunday I imported all of them, just to see."        | Music in, low. |
| 0:10–<br>0:12 | Import progress.                                                              | (silence)                                                    | Music.         |
| 0:12–<br>0:16 | Task list resolving. Count: <b>9</b> . Held.                                  | "It found nine things I'd promised people."                  | Music.         |
| 0:16–<br>0:20 | Two of the nine highlighted. Dates attached. Both in the past.                | "Three had dates on them. Two of those had already gone by." | Music dips.    |
| 0:20–<br>0:22 | <b>Hold on the two overdue items. Nothing said.</b>                           | (silence)                                                    | Music only.    |
| 0:22–<br>0:25 | LoopNote mark.                                                                | "It's LoopNote. <b>Trial link's below.</b> "                 | Music out.     |

Paid partnership · LoopNote label persists from 0:00 — platform-native disclosure, no runtime cost.

**Alt hooks:** "My phone found nine things I'd promised people." · "Oldest one I've never opened: fourteen months." **Proof:** a real backlog at 0:00; counts resolving on screen; the two overdue items held in silence at 0:20 — the viewer draws the conclusion rather than being handed it. **Objection answered:** *"I'd have to change how I work."* No — this runs on memos that already exist. One Sunday, no new habit. **Every number is a shoot-day value.** If the real count is 11, the script says 11. A real 11 beats an invented 63, and the ad's whole argument is that the evidence is real. **Predicted:** thumbstop 39–42%, trial → paid 22–24%, **CPA \$72–82.**

## Script C — "Forty seconds"

**Meta + TikTok · 9:16 · 19s** · lineage **C07** (26.8% trial → paid, best conversion in the test) + C04's moment (45% thumbstop) played straight

**Flagging the port:** C07 ran on YouTube Shorts and this brief targets Meta/TikTok. I'm porting it because it converted best in the account and its mechanism — two-state contrast — is the same one that makes C05 win on TikTok. The risk was its 32s length, so this runs at 19s.

| TIME      | ON SCREEN                                                                                           | VO / DIALOGUE                                                                                                                | SOUND                          |
|-----------|-----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|--------------------------------|
| 0:00–0:03 | Corridor, walking, phone up. <b>Phone clock legible.</b> Calendar chip: next meeting, 3 min.        | "That call ended forty seconds ago."                                                                                         | Real corridor noise. No music. |
| 0:03–0:06 | Still walking. Slightly out of breath.                                                              | "Two of the four things are already gone."                                                                                   | Corridor noise.                |
| 0:06–0:11 | Still walking. Taps record — <b>timestamp visible, matching the clock.</b> Talks fast, overlapping. | <i>(to phone)</i> "Priya needs the deck Thursday — I said I'd get pricing to Marcus, uh, and someone's chasing the vendor —" | Voice + footsteps.             |
| 0:11–0:15 | Cut to screen. Four tasks resolving. Held long enough to read all four.                             | <i>(silence)</i>                                                                                                             | UI tick. Music in.             |
| 0:15–0:17 | Task list held.                                                                                     | "I don't need a summary. I need the things I owe people."                                                                    | Music.                         |
| 0:17–0:19 | LoopNote mark.                                                                                      | " <b>Link below.</b> "                                                                                                       | Music out.                     |

**Alt hook — a planned A/B, not a backup:** "Here are the four things, thirty seconds after the call." Proof-first against the lead's dread-first. Genuine coin-flip; worth deciding with money rather than taste. **Proof:** the mess is generated live, in one take. The disorganised sentence at 0:06 and the four tasks at 0:11 are separated by five unbroken seconds, and the visible clock makes the elapsed time checkable. **Predicted:** thumbstop 35–38%, **trial → paid highest of the three at 25–27%**, CPA \$68–78.

## What I'd test next, and how I'd kill

1. **Script C's hook, dread-first vs. proof-first.** Same body, same CTA, two openings. The cleanest read available on whether this audience is moved by the cost or by the relief.
2. **Re-house C03's argument.** It converts at 22.2% inside a container nobody watches. The founder's actual reason for building this, spoken over a screen recording instead of into a lens. Cheap test, and the data says the message is already good.
3. **Deliberately break the attention band.** Every winner sat at 36–41% thumbstop. Run one hook engineered to exceed 45% and confirm it converts worse. If it doesn't, my central claim is wrong and I'd rather find out for \$500 than build a quarter on it.
4. **Kill rule:** anything over \$90 CPA at \$500 spent, no exceptions and no "give it another day." **Scale rule:** whatever holds under \$80 at \$1,500 gets the budget from the six over-target rows. Reallocating that \$21,700 at the winners' blended \$74.62 takes the account from 391 paid starts to roughly 551 — from 17% over target to 17% under — with no change to spend, offer, or landing page.

## AI usage

Required disclosure, and I'd rather over-report it.

I ran this inside **Claude Code (Anthropic)**, an agentic terminal harness, as a structured multi-pass process — data forensics, research, hook generation, drafting, adversarial critique, revision. Specifically:

- **All arithmetic** — CPA, trial → paid, cohort splits, correlations — computed in Python it wrote and I checked, not estimated. The brief's PDF was also parsed programmatically because the container had no PDF tooling.
- **Research** — web search for category positioning and the AI-notetaker trust landscape. See the limitation below.
- **Drafting and hook generation** — 30 hooks generated and cut to 8.
- **Adversarial critique** — I had it score the scripts as a hostile evaluator and it found three real defects, including cutting the best line I had written ("I wasn't behind on work. I was behind on people.") for explaining what the screen had already shown.
- **The mock LoopNote UI** — built as HTML and rendered to PNG, with the hook's "31" count verified by an assertion rather than by eye.

The judgement calls — which rows matter, what the pattern means, what to cut — are mine.

**One limitation worth stating.** I wanted verbatim user language from Reddit and App Store reviews. Reddit blocks the crawler, and the App Store, G2 and the Meta Ad Library were all blocked by the sandbox's network policy. **So I did not log a single live competitor ad, and there are no scraped user quotes in this submission** — I chose to say so rather than generate plausible ones. What I used instead was the account's own tested copy, which is arguably the better corpus: it's language that was actually put in front of this audience with thumbstop and paid conversion attached. C04's "your boss asks what was decided in the meeting" hit 45% thumbstop — the pain is validated; only the comedic framing failed it. Script C reclaims that moment straight.

**Also disclosed:** the produced version uses a **mock** LoopNote interface I built. The memo, the word count and the 31 are real; the app UI is a prop. An ad whose argument is "judge the evidence yourself" shouldn't be quiet about which part is a prop.

---

## Bonus

Script A is production-ready — mock screens built and rendered, shot list, phone settings, VO direction, edit order, export specs for 9:16 and 4:5. No talent, no location, no lighting: a screen recording and a phone mic in a parked car, about an hour.